ChatGPT

Boundary Conditions and Global Dynamics: From Holographic Principle to Attention Interfaces

Introduction

Small tweaks at a system's boundaries can have outsized effects on its overall state. In physics, information on a boundary can encode an entire volume's properties – a concept epitomized by the *holographic principle* [1] [2]. In complex networks, slight changes (e.g. a tiny increase in connectivity or a localized stimulus) can tip the whole system into a new phase, as seen in Ising models of magnetism and percolation theory's phase transitions. These theoretical frameworks offer powerful metaphors for neural and social systems, suggesting that **"boundary conditions"** – whether sensory inputs to a brain or design features at a user interface – might strongly determine *"bulk"* outcomes like global network activity or user attention patterns. In this report, we explore how boundary–bulk couplings manifest in neural networks and beyond, and how understanding these principles can inform the design of **attention modulation interfaces** (interfaces that guide or manage user attention). We will discuss the holographic principle and critical phenomena (Ising/ percolation models) and then connect these ideas to neural network dynamics (biological and artificial), social cascades, and interface design. Examples and case studies will illustrate how tweaking boundary elements can trigger large-scale changes in system behavior.

Holographic Principle: Boundary Encodes the Bulk

The holographic principle is a striking theoretical framework from quantum gravity that underscores boundary–bulk duality. It posits that **"the description of a volume of space can be thought of as encoded on a lower-dimensional boundary to the region"** [1]. In other words, all the physics inside a bounded region (the *bulk*) is fully captured by degrees of freedom on its surface (the *boundary*) [3]. Susskind famously described our three-dimensional world as *"an image of reality coded on a distant two-dimensional surface"* [4], highlighting that what happens in the bulk is mirrored at the boundary and vice versa. The prime example is the **AdS/CFT correspondence** in string theory: a 3D anti-de Sitter spacetime's gravity (bulk) is exactly dual to a 2D conformal field theory on its boundary [2]. If one changes the boundary conditions or inserts a signal at the boundary in such a duality, it manifests as a change deep in the bulk. In the AdS/CFT picture, adding energy or information on the boundary could correspond to creating a particle or perturbation propagating through the bulk. This extreme scenario illustrates how tightly coupled a boundary and bulk can be – the boundary effectively *controls* the bulk's state.

Intriguingly, analogies of holographic control appear in other domains. For instance, in machine learning theory, researchers have drawn parallels between layered deep neural networks and holographic duality. In one view, the *depth* of a deep neural network can act like an emergent extra dimension of "bulk" space, with the input layer as a "boundary" that encodes the data to be transformed [5] [6]. As data propagates inward (through hidden layers), the network is effectively constructing a bulk representation of the input, optimizing it to match the boundary conditions given by the input and output constraints [7]. In this interpretation, training the neural network – i.e. adjusting weights in deeper layers – is analogous to finding a bulk description that corresponds to the fixed boundary data. While largely metaphorical, this AdS/Deep

Learning correspondence suggests that **small changes at the input boundary can fundamentally alter the learned internal representation** – much as a tweak in boundary fields in AdS/CFT would change the bulk geometry or fields [5] [6]. The holographic principle thus provides a conceptual framework where *interfaces* (boundaries) and *states* (bulk configurations) are two sides of the same coin. It encourages us to look at system interfaces (literal or abstract) as potent loci of control over the system's interior dynamics.

Critical Phenomena: Ising Models, Percolation, and Phase Transitions

Another set of frameworks highlighting boundary-triggered global changes comes from **critical phenomena** in statistical physics – notably the Ising model of magnetism and percolation theory. These models illustrate how **small local changes can lead to** emergent global order**, especially when a system is tuned near a critical threshold.

Ising Model: The Ising model consists of many interacting "spins" on a lattice that can be +1 or –1, akin to tiny magnets [8]. Each spin tends to align with its neighbors, and the competition between this coupling and thermal randomness produces a phase transition. At high temperature, spins are disordered (no net magnetization), but below a critical temperature *T_c* they spontaneously align, magnetizing the entire lattice. Near *T_c*, the system is exquisitely sensitive to small perturbations: **an infinitesimal external magnetic field or a slight bias in boundary spins can "flip" the entire system's magnetization** from majority-up to majority-down. This is because the correlation length at criticality becomes very large – spins separated by long distances become correlated – so a local fluctuation can propagate its influence system-wide. In essence, the **boundary conditions in an Ising system (e.g. fixed spins at the edges) can determine the bulk state** when the system is in or near the ordered phase. For example, if you "seed" the boundaries of an Ising lattice with spins pointing up while the system is slightly below *T_c*, the bulk will tend to align up as well, producing a globally magnetized state. Once the system is in the ordered phase, it has two possible stable global states (all spins mainly up or mainly down), and a small nudge at the boundaries or a tiny global field can tip the magnetization from one to the other – a dramatic non-linear response to a minor input. This threshold behavior is a hallmark of symmetry-breaking phase transitions.

Notably, **brains and neural networks may themselves operate near such critical points**. Empirical studies of human brain activity at "rest" (no explicit task) have found patterns consistent with an Ising model at critical temperature [9]. For example, one study compared functional MRI brain networks to 2D Ising simulations and found *"striking similarities"* at *T_c*, suggesting the brain may tune itself to a critical state between order and disorder [10]. In a critical brain, small neural fluctuations don't just fade away; instead they can cascade into system-wide events (sometimes called neuronal avalanches). This means a tiny stimulus or a slight change in a single cortical area could evoke a disproportionately large activation pattern across the whole brain if the network is poised at criticality. In other words, the brain might leverage a physics principle – hovering at the boundary between phases – to maximize responsiveness and range of influence for minimal inputs. It's an attractive notion that connects to theories of attention and consciousness: the brain could rapidly reconfigure global activity (bulk state) in response to slight "boundary" inputs (sensory cues or internal signals) because it's at a sensitive critical point.

Percolation Theory: Percolation is another paradigm for **small changes causing a sudden global connectivity**. In percolation models, we occupy sites or links in a network with some probability *p*, and ask whether a giant connected cluster spans the system [11] [12]. There is a critical probability *p_c* at which a phase transition occurs: below *p_c* the network breaks into small clusters, but above *p_c* a **spanning cluster** forms that connects across the entire system [13]. At *p = p_c*, the addition or removal of just a few links can be decisive – one extra bond might suddenly connect two large clusters, **joining them into one giant component that percolates from one boundary of the system to the other** [14] [13]. This is a geometric form of a phase transition. It implies a kind of *tipping point*: when a network is just at the critical connectivity, **a minor "boundary tweak" – e.g. opening one more channel or connection – can yield a system-wide change from fragmentation to integration**. Conversely, closing a single critical connection can shatter a formerly whole network into pieces. Percolation concepts have been applied widely outside of materials science, including to **epidemics and social contagion**. For instance, the spread of a virus or a piece of information on a social graph can exhibit a percolation threshold: below a critical density of contacts the outbreak dies out locally, but beyond that threshold a giant outbreak or information cascade ensues, infecting a finite fraction of the entire network [15] [16]. Duncan Watts's threshold model of social cascades is a classic example: if an initial seed group exceeding some critical size or fraction is "activated" with a new idea, a large cascade can be triggered across the network [17] [18]. This model shows a **phase transition** in cascade size – either a small isolated adoption or a massive cascade – depending sensitively on the network connectivity and the initial seed fraction [15] [19]. In sparse networks, cascade sizes follow a broad distribution (some cascades die out, a few go large), whereas in denser networks the distribution becomes bimodal (either negligible or system-spanning) [20]. The key point is that **the boundary between a tiny ripple and a global wave often hinges on hitting a critical threshold**. At that threshold, a *boundary modification* (like adding one more early adopter, or one more link between communities) can cause a qualitative shift to a new global state (a widespread adoption cascade) [18].

These critical phenomena teach an important lesson: **when a system sits at a delicate balance point (criticality), the leverage of small inputs is maximal**. The system's internal correlations or connectivity are so extensive that a perturbation at one boundary can percolate throughout. In contrast, when a system is far from critical (e.g. far below percolation threshold, or far above *T_c* in the disordered phase), boundary tweaks tend to remain local, having only limited, proportional effects. Thus, to achieve large-scale modulation via boundary changes, the system should be engineered or maintained near a critical point or phase transition zone.

Neural Networks: Boundary Inputs and Global Outputs

Focus first on **neural networks**, both artificial and biological, as these are central to cognitive science and attention. Neural systems often exhibit complex, non-linear dynamics where small changes in input or connectivity can yield big changes in activity patterns – especially if the network is in a high-gain or critical regime.

Artificial Neural Networks and Adversarial Examples

In artificial deep neural networks (DNNs), the *input layer* can be viewed as the "boundary" where data enters, and the network's final output is a result of many layers of internal transformations (the "bulk" computation). It has been dramatically demonstrated that **minute changes at the input can completely alter the network's global decision**, even if those changes are imperceptible to humans. A striking case is that of **adversarial examples**: by adding a tiny, carefully crafted perturbation to a normal input (such as an image), one can fool a state-of-the-art classifier into a wrong prediction with high confidence [21].

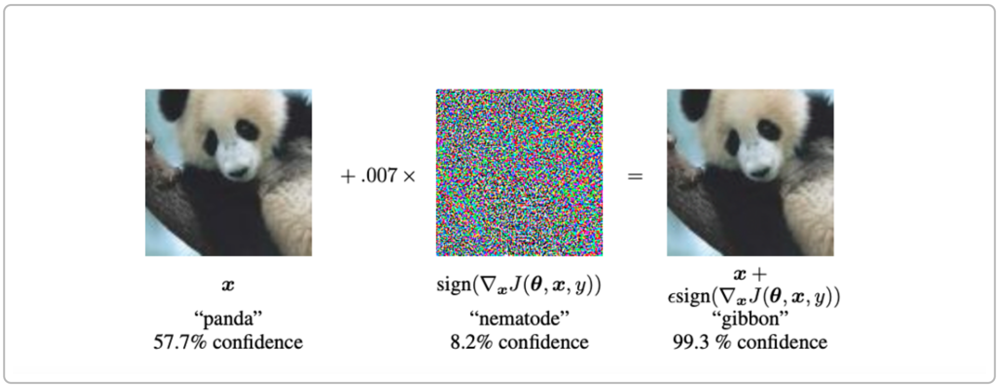

An example of an adversarial input perturbation. A convolutional neural network originally classifies the left image as a "panda" with 57–60% confidence. Adding an imperceptible noise pattern (middle) to the image yields the right image, which the network confidently (≈99% confidence) misclassifies as a "gibbon." The network's high-level output (bulk behavior) is thus flipped by a virtually invisible boundary tweak [21] *.*

This famous panda-versus-gibbon example (from Goodfellow *et al.*, 2015) illustrates the extreme sensitivity of deep networks: the decision boundary in input space can be very complex, allowing a tiny input vector change to push the example across into a totally different category [21] . In effect, the *entire* internal activation pattern of the network – and hence the final classification – was changed by a perturbation of only 0.007 in pixel values [22] . Such adversarial vulnerability is analogous to a physical system at criticality: the network was right at a boundary (of its classification manifold) where a minute nudging yields a large effect (class A → class B). It underscores that **for complex learned models, "boundary conditions" (input signals) can exert disproportionate control over "bulk" computations**. Designers of AI systems thus must be mindful that a system can be thrown into a wholly different state by what looks like a negligible input change – a concept resonant with chaos theory as well.

Beyond adversarial examples, the architecture of deep networks itself hints at a boundary-bulk relationship. Each hidden layer transforms the representation of data, analogous to successive coarse-graining or feature extraction – reminiscent of renormalization group (RG) flows in physics [23] . In fact, theoretical work has drawn correspondences between deep networks, RG, and even holographic mappings [24] [7] . The idea is that the input layer encodes information which is then "projected" through many nonlinear transformations (the bulk), yielding an abstracted output at the final layer. If we view the input as defining a boundary condition (a particular configuration entering the network), the propagation of activation through layers is like the diffusion of that boundary influence inward. Small differences in the input can get amplified or reinterpreted by the time they reach deep layers, leading to divergent outputs. This is routinely observed in pattern recognition networks: e.g. a few pixel changes can cause one high-level category to be activated over another, because those pixels might flip the sign of some feature in later layers. **Energy-based neural networks** such as Hopfield networks and Boltzmann machines make the link to physics even more explicit – they are essentially equivalent to Ising spin systems searching for low-energy states (stored memories) [23] . In a Hopfield network (which can serve as an associative memory), giving the network a partial or noisy pattern is like setting certain neurons (boundary) and letting the others evolve – the network will settle into the nearest stored pattern (a global attractor). A small change in the initial cue can lead the network to converge to a completely different memory, if that change nudges it closer to another attractor's basin. This again is a case where *tweaking initial conditions yields a wholesale shift in system outcome*. In summary, artificial neural nets demonstrate both through adversarial phenomena and theoretical analogies that controlling boundary inputs is a powerful way to control the internal state – and that these systems can exhibit tipping-point behavior where minimal input differences yield major output differences.

Brain Networks, Criticality, and Attention

Biological neural networks (brains) similarly exhibit nonlinear dynamics where boundary inputs (sensory stimuli, or external perturbations) can have global impacts – especially if the brain is in a poised, sensitive state. As noted, evidence is mounting that the brain operates near criticality [10]. At this edge-of-chaos state, neurons form dynamic correlation networks with power-law distributed activity cascades (neuronal avalanches). In practical terms, this means the brain can **integrate small inputs over large scales** – a tiny stimulus can recruit many neurons into synchronous activity if conditions are right. For example, a faint but meaningful cue in the periphery of your vision might suddenly grab your attention and trigger a widespread shift in cortical activity (as anyone who has seen a spider out of the corner of their eye can attest!). This is because the brain's resting-state networks are teetering on a connectivity threshold where functional clusters can rapidly expand [25]. If a stimulus activates neurons that belong to a strongly interconnected cluster (a *percolating cluster* in cognitive terms), the activation will **percolate** through that cluster and influence a large population of neurons [25]. On the other hand, stimuli that hit isolated circuits (non-percolating regions) might fizzle out with little global effect. Thus, the brain's effective connectivity (which pathways are "open" or "closed" to signal flow) determines whether a local event stays local or causes a global broadcast.

Importantly, **attention** is one of the brain's mechanisms for dynamically tuning these "percolation pathways" in real time. Attentional modulation literally changes effective connectivity and gain in neural circuits [26]. When you focus on a particular sensory input or task, your brain boosts the synaptic weights and synchrony in circuits handling that information, while suppressing distractions – in effect **opening the spigot for certain signals and closing it for others** [26]. Neuroscience studies show that attention increases the firing rates of neurons representing attended stimuli, reduces noise correlations, and can even alter network oscillatory states [27] [28]. In terms of our models, attention moves the brain closer to a phase where signals in the attended pathway form a large, resonant cluster of activity (so that any input in that pathway has a big downstream effect), whereas unattended inputs face an "insulating" network that prevents their spread. One could say the brain **modulates its own boundary conditions**: by internally gating inputs (top-down attention), it decides which sensory channels are effectively allowed to influence the global workspace of the mind. This is analogous to changing the open bond probability p in percolation – attention might raise p for neurons in the spotlight (strengthening their connections) and lower p elsewhere, thus ensuring that an attended signal will percolate across cortex and dominate awareness.

It's noteworthy that **emotion and other top-down factors** also play a role in this; they can bias attention and information flow (for instance, anxiety might make even small threats highly salient, effectively increasing their "weight" in the network). The brain's neuromodulatory systems (acetylcholine, norepinephrine, etc.) can shift it into different global states (e.g. a vigilant state vs. relaxed state), altering how readily local events propagate. In a vigilant state (high arousal), the cortex might be closer to criticality – ready to amplify any anomaly – whereas in a drowsy state it might be subcritical, with inputs less likely to cause widespread shifts. Thus, through evolution the brain might have learned to toggle these conditions to balance sensitivity and stability. When we need to concentrate, we intentionally isolate ourselves from other inputs (closing boundary channels like silencing your phone) to avoid unwanted cascades of distraction. Conversely, when scanning for danger or opportunity, we might *increase* our receptivity to peripheral cues.

In summary, neural networks – whether engineered or natural – exemplify boundary→bulk coupling: initial conditions or boundary inputs can select which attractor or global state the network settles into. Moreover, these networks can be tuned to a knife-edge (critical point) where such selection is especially potent. This has direct implications for designing interfaces that interact with human neural systems, as we explore next.

Social Networks and Cascades

While our focus is on neural systems, it's worth noting that **social and information networks** also display boundary-triggered global phenomena. A society or online network can be seen as a vast system of nodes (people) with connections (communication links). Here, the "boundary" might be a subset of influential individuals or an initial seed of information injected into the network. We often observe that *small groups or single actors at the edge of mainstream discourse can spark large collective changes*, if the social network is susceptible. For example, a single whistleblower or a viral tweet (a tiny boundary event in the grand scheme) can ignite a worldwide social movement or a rapid mass opinion shift. This resembles the threshold and percolation behavior discussed earlier.

In the **global cascades model** of complex contagion, there is a critical threshold condition for widespread adoption [17] [18]. If an initial "activated" fraction or an influential cluster in the network exceeds that threshold, a domino effect occurs and a large fraction of the network ultimately adopts the new idea or behavior [18]. Below the threshold, the innovation fizzles out without systemic impact. Empirical studies confirm such dynamics: for instance, one analysis of collaboration networks found that **indirect influence** (friends-of-friends, etc.) can *dominate* direct influence in spreading behavior, meaning that network structure enables small seeds to have cascading reach beyond their immediate neighbors [29]. This is essentially percolation in action – a person convinces their friends, who convince their friends' friends, and so on, creating a large cluster of influence that wasn't obvious from the local start. The connectivity and degree distribution of the social network are crucial in determining this tipping point [15]. Highly connected hubs can act as powerful boundary points; if a hub converts, it can rapidly transmit information to many parts of the network (in percolation terms, hubs help form the spanning cluster) [15] [19]. On the other hand, networks that are more modular or fragmented may resist global cascades – influences stay trapped in local clusters unless a *bridge* (boundary spanner) exists to carry them across communities.

A real-world example is the spread of new technologies or behaviors: sometimes a niche group (early adopters) remains isolated (like a cluster that doesn't percolate), but if one of those early adopters has connections into the wider population (a boundary link), suddenly the innovation jumps and becomes mainstream. The concept of the "tipping point" in social phenomena, popularized by Malcolm Gladwell, aligns with these network phase transitions: a small input at just the right point, when the system is ripe, unleashes a big change. Conversely, to prevent undesirable cascades (e.g. misinformation or epidemics), interventions often target boundary nodes – for instance, **vaccinating or quarantining a few well-connected individuals can halt an epidemic's spread by breaking key links**, just as removing a single critical bond can prevent a spanning cluster in percolation.

The lesson from social network dynamics is that structure and initial conditions matter enormously: one cannot assume proportional scaling (10× the influence leads to only 10× the effect); rather, there are nonlinear thresholds. For designers of communication systems or social platforms, this means a small design change at the "boundary" of user interaction might cascade – for good or ill – through user behavior. Features like retweet buttons, algorithmic feed tweaks, or group recommendations are seemingly minor interface details, but they set boundary conditions that can foster echo chambers or viral cascades globally. Recognizing these effects can help in **modulating the spread of information and attention in social systems** more deliberately.

Attention Modulation Interfaces: Applying Boundary–Bulk Principles

Finally, we turn to **practical applications in interface and product design**, especially regarding **attention modulation**. An "attention modulation interface" refers to any system (software, device, or UI/UX feature) that aims to guide, capture, or manage where a user's attention goes. Given the theoretical insights above, treating the interface as the *boundary* to the user's cognitive system provides a powerful design perspective. By adjusting interface elements, we effectively tweak the boundary conditions of the user's perceptual and cognitive inputs – which can lead to large changes in their mental state or behavior (the bulk response).

Modern users are immersed in an environment of competing stimuli: *"notifications, banners, competing content, and infinite scrolls constantly vying for a person's focus"* [30] . These interface elements are boundary inputs to our brain's attention network. Poorly managed, they can induce a chaotic, distracted state in the user – analogous to pushing the brain toward the disordered phase where no single signal dominates. Thoughtful design, however, can do the opposite: it can raise certain signals above the "noise floor" and suppress irrelevant stimuli, effectively moving the user's cognitive state into a more ordered, focused regime. In terms of our earlier metaphors, a good interface can *tune the percolation parameter p of information flow*, ensuring that task-relevant information "percolates" through the user's mind while distractions remain isolated.

Case Study – Notification Design: Consider the simple act of receiving notifications on a smartphone. If an app bombards the user with frequent pings (high p of distractions), even minor messages can hijack global attention – the user's train of thought is interrupted and their focus shifts to the phone. This is a boundary tweak (a tiny sound or icon) causing a bulk change (the user stopping work to check a message). In contrast, *focus mode* features or Do-Not-Disturb settings alter this boundary condition: they batch or silence notifications, effectively "closing" those input channels. Users often report significant improvements in concentration and stress when they implement such boundary controls. Thus a relatively small interface change (e.g. changing a default notification setting from opt-out to opt-in) can yield a large behavioral difference across the user base, as many more people remain in sustained attention. From a product perspective, this is boundary–bulk coupling: the interface boundary was tweaked to modulate global user engagement patterns.

Case Study – Peripheral Cues for Attention Guidance: Another example is using subtle peripheral cues to guide attention within an interface. Research in human-computer interaction suggests that peripheral visual cues (like a gentle highlight or a movement in the corner of the screen) can draw a user's gaze without overly disrupting them, whereas a full pop-up or modal dialog seizes focal attention abruptly. The peripheral cue acts as a small boundary nudge – if timed and placed well, it can gradually redirect the user's "bulk" cognitive focus to a new area or task. For instance, an interface might *blink an icon softly* to indicate something needs attention; the user, if at a receptive moment, will shift attention and address it. This is reminiscent of initiating a phase transition in the brain's attention state: a minimal cue, if the user is near a decision threshold, triggers a full context switch (where their mental resources reallocate globally to the new task). Such techniques leverage our brain's sensitivity to change in the periphery – a trait developed for survival – but doing so in a controlled, non-intrusive way uses the boundary→bulk idea elegantly. The user interface *boundary* provides just enough input to tip the *bulk* cognitive state in the desired direction.

Case Study – Cognitive Brain-Machine Interfaces: At the high-tech end, we have *closed-loop brain-machine interfaces* that explicitly measure a user's neural state and adjust stimuli in real time. A recent example is a cognitive brain-machine interface (cBMI) used to track and enhance attention [31]. In one study, researchers monitored a specific EEG signal (the steady-state visually evoked potential, SSVEP) as an index of a person's attention level [32] [33]. The interface was programmed to **intervene with a task stimulus only when attention was high or low**, to test and train the effects of attention fluctuations [34]. Essentially, the system closes the loop: the brain's current "bulk" state (engaged vs. disengaged) is measured at the boundary (scalp EEG), and the interface then provides feedback or new input to modulate that state [31]. For example, if the user's attention lapses (SSVEP power falls), the interface might introduce a novel stimulus or audio alert to bring them back (a boundary jolt to push the brain back toward the attentive phase). Conversely, if attention is high, it might present a challenge or reward to reinforce that state. Such cBMI setups demonstrate in a very literal sense how boundary signals can drive bulk neural dynamics: *"a cognitive BMI provides closed-loop neurofeedback based on neural signatures of cognitive processes, like attention"* [31]. Early results show that users can be trained with these feedback loops to maintain better focus, essentially learning to self-regulate into a desired global state when prompted. This mirrors the principle of gentle nudging at the boundary to cause a beneficial whole-brain outcome.

Design Implications: The above cases suggest a few design principles for attention-modulating systems: (1) *Operate near the user's cognitive critical point*: A good interface neither under-stimulates nor overloads, but keeps the user in a zone where they are alert and inputs have impact. In this state, well-timed small cues can have maximal effect (just as small signals propagate far at criticality). (2) *Use minimal effective interventions*: Aim for the smallest boundary tweak that achieves the desired bulk change. This is both ethical (avoids heavy-handed manipulation) and effective (users are less likely to develop fatigue or reactance). For example, rather than a loud alarm when productivity wanes, an app might subtly change the color tone of the screen or play a gentle sound – enough to nudge the user back on track without a jarring disruption. (3) *Personalize the threshold*: Different individuals have different "phase transition" points for attention and arousal. Adaptive systems could learn an individual's responsiveness and adjust the intensity of boundary cues accordingly, much like tuning p to just the point of percolation for each user. (4) *Beware unintended cascades*: Just as in social networks a small design change can lead to viral misinformation, in cognitive interfaces a small tweak might have side effects. Designers should simulate and user-test for cases where an innocent nudge spirals into over-stimulation or distraction. For instance, a notification system that tries to encourage engagement could accidentally push a user into a multi-hour distraction (a cascade of clicks) – essentially a design-induced attention avalanche.

In a product like **"the Donut"** (if we assume it's an attention-related tool), these concepts might be concretely applied. The Donut could, for example, present a ring-shaped visualization of focus (boundary) that tightens or relaxes based on user's current attention level (bulk state), prompting the user to adjust. Or it might create a *holographic* effect where the user's peripheral vision (boundary of the visual field) subtly mirrors their central task performance – providing an intuitive cue encoded at the edges that reflects the whole. The holographic principle's influence could be metaphorical: designing the interface such that *"the state of the whole task is visible at a glance in one element"* (like a boundary summary that encodes the bulk progress). Likewise, using Ising/percolation ideas, the Donut's UX might incorporate a critical threshold – for example, a timer that turns red only when distractions exceed a certain rate, causing the user to suddenly notice and correct course (a phase transition from “flow” to “distracted” signaled by a boundary indicator, which then spurs a global refocus effort).

Conclusion

Across theoretical physics, network science, and interface design, we see a unifying theme: **what happens at the boundary can determine the fate of the bulk**. The holographic principle gives the ultimate image of this unity, treating the boundary and bulk as equivalent in content. Critical phenomena like those in Ising and percolation models show that systems can be exceptionally sensitive to boundary or initial tweaks when poised at transition points. Neural and social systems, with their myriad interactions, often inhabit these transition zones, making them capable of both stability and rapid phase change. For designers of interactive systems and *attention modulation interfaces*, these insights are invaluable. They remind us that small, well-crafted changes to the information presented to users (the interface boundary) can lead to disproportionate improvements in user experience and behavior – such as enhanced focus, better information spread, or more ethical engagement – if applied at the right time and context. The boundary–bulk coupling also cautions us: manipulative or careless boundary inputs can likewise drive negative global outcomes (chaos, distraction, viral harm). Thus, a deep understanding of boundary effects enables us to **engineer interfaces that harness positive cascades and avoid destructive ones**.

In summary, whether we are encoding an entire universe on a 2D surface or simply trying to get a user to pay attention to a critical alert, the principle is the same: *tune the boundaries thoughtfully, and the desired bulk state will follow*. By learning from physics and network theory, we can design smarter systems that use minimal intervention for maximum beneficial impact – guiding complex brains and societies through gentle nudges rather than forceful shoves, and creating order from the edge.

Sources:

- Susskind, L. (1995). *The World as a Hologram.* (Quote on the world as a hologram) [1] [2]
- Bousso, R. (2002). *The holographic principle.* (General description of physics being encoded on a boundary) [3]
- Ganchev, A. (2019). *On bulk/boundary duality and deep networks.* (Deep NN layers as emergent bulk dimensions) [5] [6]
- Fraiman, D. et al. (2009). *Ising-like dynamics in large-scale functional brain networks. Phys. Rev. E*, 79(6). (Brain resting-state near critical Ising point) [9]
- *Percolation theory – Wikipedia:* Explanation of percolation threshold and phase transition [13] [14]
- Watts, D. (2002). *A simple model of global cascades on random networks. PNAS*, 99(9). (Threshold fraction for cascade, phase transition in cascade size) [17] [18]
- Indirect influence and percolation in social networks (PNAS 2021 study) [29]
- *Adversarial example (panda vs gibbon) – Goodfellow et al. (2015) via KDnuggets:* Demonstration of tiny input noise causing misclassification [21]
- Goodfellow, I. (2015). *Explaining and Harnessing Adversarial Examples.* (Original paper for adversarial panda example) [21]
- ScienceDirect Topics: *Attentional Modulation.* (Definition: attention influences connectivity between cortical regions) [26]
- Premereur, E. et al. (2011). *Attentional modulation of neural processing.* (Attention alters firing rates, correlations in visual cortex) [27] [28]

- Joon Kim, S. et al. (2022). *Tracking fluctuations in attention with a cognitive brain–machine interface. Comm. Biology*, 5, 1328. (Closed-loop EEG-based attention feedback system) [31] [34]
- Curtis, M. (2024). *UX and the Psychology of Attention. Medium.* (Challenges of capturing focus ethically in UX design) [30]

---

[1] [2] [3] [4] Holographic principle - Wikipedia
https://en.wikipedia.org/wiki/Holographic_principle

[5] [6] [7] arxiv.org
https://arxiv.org/pdf/2401.00939

[8] Ising model - Wikipedia
https://en.wikipedia.org/wiki/Ising_model

[9] [10] [0811.3721] Ising-like dynamics in large-scale functional brain networks
https://arxiv.org/abs/0811.3721

[11] [12] [13] [14] Percolation theory - Wikipedia
https://en.wikipedia.org/wiki/Percolation_theory

[15] [16] [17] [18] [19] [20] Global cascades model - Wikipedia
https://en.wikipedia.org/wiki/Global_cascades_model

[21] [22] Deep Learning Adversarial Examples – Clarifying Misconceptions - KDnuggets
https://www.kdnuggets.com/2015/07/deep-learning-adversarial-examples-misconceptions.html

[23] [24] The Mathematics of the Ising Model and Neural Networks by Miquel Noguer I Alonso, Can You :: SSRN
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=5078967

[25] Percolation (cognitive psychology) - Wikipedia
https://en.wikipedia.org/wiki/Percolation_(cognitive_psychology)

[26] Attentional Modulation - an overview | ScienceDirect Topics
https://www.sciencedirect.com/topics/engineering/attentional-modulation

[27] [28] [31] [32] [33] [34] Tracking momentary fluctuations in human attention with a cognitive brain-machine interface | Communications Biology
https://www.nature.com/articles/s42003-022-04231-w?error=cookies_not_supported&code=fd10ea1c-b30b-4584-b6ea-7861a8f953ae

[29] Indirect influence in social networks as an induced percolation ...
https://www.pnas.org/doi/10.1073/pnas.2100151119

[30] UX and the Psychology of Attention | by Mike Curtis | The You Design System | Medium
https://medium.com/mikewcurtis/ux-and-the-psychology-of-attention-9f050e192328