



Phase Oscillators for Control, Coordination, and Attention Routing in AI Systems

Introduction

Oscillatory dynamics – where components cycle through phases over time – offer powerful tools for organizing behavior in complex systems. Recent research has leveraged **phase oscillators** (rhythmic signal generators) to improve control and decision-making in AI. From **robotics and reinforcement learning (RL)**, to **multi-agent coordination**, and even **neuroscience-inspired cognitive architectures**, phase-based mechanisms enable robust performance and flexible information routing. Below, we survey how oscillators are used in these domains, highlighting both theoretical foundations and practical implementations, and discuss what behavioral benefits (e.g. robustness, efficiency) they bring.

Oscillator-Based Control in Robotics and Reinforcement Learning

*Figure: A Unitree A1 quadruped robot carrying a 13.75 kg payload (plates on its back) using a phase-oscillator controlled gait. The oscillator-driven controller allows the robot to maintain stable locomotion even with this heavy load, demonstrating the robustness imparted by a CPG-RL (Central Pattern Generator + Reinforcement Learning) approach 1. Central Pattern Generators (CPGs) are networks of coupled neural oscillators that produce rhythmic motor patterns. Controllers based on CPGs exploit **entrainment** – the tendency of coupled rhythms to synchronize – which makes motions **robust against perturbations**. For example, if a disturbance causes a leg to land off-timing, the oscillators quickly re-lock phase and the gait recovers its rhythm 2. By embedding such oscillatory circuits into a robot's control policy, robots gain a form of adaptive “self-stabilization” that traditional rigid controllers lack.*

In practice, combining **deep RL with CPG controllers** has proven effective for locomotion. Instead of learning gaits entirely from scratch, RL can **tune the parameters of an oscillator network** that drives the robot's joints. Notably, a CPG is characterized by a few key features – the phase timing between limbs, the shape of the motion waveform, and the oscillation frequency 3. Recent work has used simplified phase-oscillator models where these features correspond directly to tunable parameters, and then applied RL to optimize those parameters for different terrains 4. This approach narrows the search space (injecting prior knowledge of rhythmic gait structure) and yields fast adaptation. For instance, Sakai *et al.* (2025) report that their quadruped agent could adjust to new surface conditions within ~1,500 trials by learning to modify oscillator phase and frequency, achieving suitable gaits in only \$5\times 10^5\$ time-steps 5 6.

Critically, **RL agents can learn to modulate oscillator outputs** rather than raw joint torques. Bellegrada and Ijspeert (2022) demonstrated a deep RL policy that *directly adjusts the CPG oscillators' amplitude and frequency setpoints*, coordinating a quadruped's four legs via phase-aligned signals 7. The resulting controller inherently generates smooth, periodic motions but can also adapt them based on feedback. In simulation-to-reality tests, this oscillator-based policy proved remarkably **robust** – the real robot could tolerate a dynamically added payload **115% of its body mass** (see figure) and traverse uneven terrain without special tuning 1. Notably, the structured CPG approach reduced the need for extensive domain

randomization or sensing; even with minimal sensing (just binary foot contact signals), the learned policy remained stable ¹. Such resilience stems from the oscillators' entrained dynamics and coupling: the agent's actions are **temporally coordinated** and can absorb shocks or delays by phase-adjusting rather than breaking sequence. Overall, incorporating phase oscillators into RL controllers has yielded **more sample-efficient learning and robust behavior** than purely feedforward neural policies, especially in rhythmic tasks like walking. Researchers are now exploring how this paradigm can illuminate biology (e.g. understanding the role of descending pathways and reflex feedback in gait by observing which oscillator parameters the RL tweaks) ⁷ ⁸, further bridging control theory and neuroscience.

Phase Synchronization for Multi-Agent Coordination

In multi-agent systems, **synchronization through oscillators** provides a natural mechanism for coordination. In fact, the coordinated behavior of distributed agents is often compared to oscillator synchronization in physics ⁹. Many phenomena in nature illustrate how coupled oscillators can self-organize a group: for example, pacemaker cells in the heart sync their firing to beat in unison, and colonies of fireflies will even **lock their flashing phases** when enough of them are coupled visually ¹⁰. These examples show that individual units, by following simple rules of phase coupling, can spontaneously “agree” on a common rhythm to achieve a collective goal (whether it’s pumping blood or creating a unified display). This insight inspires engineered multi-agent systems to use phase alignment as a **decentralized coordination strategy**.

Recent research indeed adapts mathematical oscillator models to multi-agent AI. Mitra (2025) proposes modeling each AI agent as an **oscillator with a phase** (and possibly amplitude) that evolves over time and is influenced by other agents ¹¹. By using a classic Kuramoto-style coupling (a sinusoidal attraction of phases), the agents gradually synchronize their internal rhythms. An **order parameter** can quantify the group's level of coordination – essentially measuring if all agents are in phase or dispersed ¹². Using this framework, one can tune parameters like coupling strength or network connectivity and study their effect on emergent teamwork. Simulations have shown that **stronger coupling leads to robust synchronization**: even if the agents have heterogeneous abilities or dynamics, increasing their mutual influence causes the phases to align and the group to behave coherently ¹³. In Mitra's results, when oscillators were tightly coupled, the multi-agent system achieved high phase synchrony (high order parameter) and could maintain it despite differences between agents ¹⁴. This suggests that by simply adjusting interaction gains, one can drive a team from disorganized to **unified timing**, improving cooperation on tasks that require simultaneity or sequencing.

Beyond pure synchronization, phase-based models offer insight into **multi-agent decision-making processes**. Interestingly, the oscillator framework has been linked to *chain-of-thought reasoning*: iterative problem-solving by a team of agents can be seen as a sequence of state updates that benefit from being in sync ¹⁵. If agents (or different “thought” steps) share a rhythm, they may exchange information more effectively – analogous to how humans brainstorming together often fall into a synchronized pace. The Kuramoto-inspired model formalizes this by showing a correspondence between synchronized phase dynamics and coherent group reasoning ¹⁶. This **physics-informed approach** provides a rigorous way to design and analyze multi-agent AI. Instead of relying purely on heuristic communication protocols, we can leverage oscillator coupling theory to ensure scalability and stability: the system will naturally converge to coordinated behavior if designed with the right coupling parameters. Overall, treating agents as oscillators opens a path to **adaptive and interpretable coordination** strategies – for example, one might adjust coupling in real-time to flexibly regroup agents or to split a team into sub-synchronized clusters for

multitasking. Such ideas are early-stage but highlight the promise of oscillatory control principles for complex multi-agent orchestration ¹⁴.

Oscillatory Mechanisms in Cognitive Architectures and Attention Routing

Neuroscience has long suggested that the brain uses rhythmic oscillations to **route information and govern attention**. Different frequency bands of neural oscillation are associated with different functional roles, providing a multi-scale clock for brain activity ¹⁷. For instance, theta oscillations (~4–8 Hz) in the hippocampus facilitate memory encoding and navigation, alpha rhythms (~8–12 Hz) in visual areas reflect an idling or inhibitory state (often linked to suppressing distractions), beta (~13–30 Hz) involves active cognitive processing and motor planning, and fast gamma oscillations (~30–100 Hz) are known to **coordinate high-level perception and attention** ¹⁷ ¹⁸. This hierarchy of rhythms is thought to organize brain computations: slower waves set broad timing for large-scale integration, while faster oscillations handle local binding and feature grouping. A concrete example is **selective visual attention**. Experiments show that when an animal attends to a particular stimulus, neurons encoding that stimulus synchronize their firing with a local gamma oscillation, whereas neurons processing ignored stimuli fall out of sync ¹⁹. Because gamma oscillations create a rhythmic fluctuation in neuronal excitability, spikes arriving “in phase” (at the high-excitability part of the cycle) are more likely to propagate, while out-of-phase inputs are dampened ¹⁹. In effect, oscillatory phase acts as a **gating mechanism**: only inputs aligned to the correct phase are passed along, thereby *routing attention* to the chosen stimulus and suppressing others. Theoretical models have demonstrated that such **phase-dependent gating** can select among competing input streams, and empirical measures show that attention increases gamma-phase coherence between brain regions handling the attended information ²⁰ ²¹. This principle is often dubbed “*communication through coherence*” – groups of neurons (or brain areas) that oscillate in unison can effectively communicate, whereas those out of sync are functionally disconnected ²⁰. It provides a bio-inspired blueprint for **attention routing**: by controlling the timing (phase) of signals, a system can flexibly enable or block pathways without changing the physical connections.

Motivated by these insights, AI researchers are injecting oscillatory dynamics into artificial neural networks to improve temporal coordination and attention handling. A straightforward example is the **Phased LSTM** (Neil *et al.*, 2016), a variant of the Long Short-Term Memory network that includes a learned cyclical gate. This *time gate* is controlled by a parametrized internal oscillation, which means each LSTM unit opens to process inputs only during specific phase intervals and remains closed otherwise ²². By updating the network’s state in rhythmic bursts, the Phased LSTM can naturally handle irregular or event-driven data – akin to a neuron that samples information in discrete pulses – and showed faster training on tasks with sparse time structure. Essentially, it gave an RNN a sense of “clock ticks” to **schedule its attention** to incoming events.

More recent architectures incorporate oscillators more deeply into network structure. Rajagopal *et al.* (2025) introduced a **Deep Oscillatory Neural Network (DONN)**, where certain neurons are replaced with dynamical elements that behave like Hopf oscillators (producing sinusoidal activity in the complex domain) ²³. The DONN combines these oscillatory neurons with standard neurons (using complex-valued weights) and is trained via backpropagation (extended to complex numbers). Impressively, this network not only matched or exceeded the performance of conventional networks on signal processing and image recognition benchmarks, but it also exhibited emergent **brain-like behaviors**. For example, during image

classification, the network's oscillatory neurons tended to synchronize in patterns that achieved **feature binding** – grouping visual features that belong to the same object via phase coherence, reminiscent of how the visual cortex binds features (color, shape, motion) of an object by synchronizing the neurons representing those features ²⁴. Additionally, when trained with a Hebbian (unsupervised) learning rule, the network naturally manifested a spike-timing-dependent plasticity (STDP) curve in its weight updates ²⁴, indicating that the oscillatory timing was leading it to a biologically plausible learning behavior. These phenomena suggest that adding oscillators to deep networks can make internal representations more **temporally structured and interpretable** – the timing of activation becomes meaningful, not just the magnitude, enabling the network to internally route information by phase just as the brain might.

Another line of work by Wang *et al.* (2025) takes inspiration from multi-frequency brain organization. They proposed a **tripartite neuromorphic architecture** with three interacting modules (for sensory processing, modulatory context, and executive decision-making) and introduced **multi-frequency oscillatory signals** to synchronize these modules ²⁵. Essentially, different subsystems operate with their own rhythmic activity, but a coordinating mechanism aligns their phases appropriately across **multiple timescales** (e.g. a slower oscillation linking high-level decisions and a faster oscillation linking low-level feature detectors). This design led to notable improvements: on visual tasks, the oscillation-enhanced architecture achieved about *2.18% higher accuracy* than a comparable non-oscillatory model and did so with *48% fewer processing iterations*, indicating more efficient information flow ²⁶. Moreover, the model's outputs correlated more strongly with human confidence ratings, suggesting that the oscillatory synchronization endowed it with more human-like dynamics of decision certainty ²⁶. The authors attribute these gains to the **temporal coordination across specialized modules** – by imitating how the brain's theta, beta, gamma, etc. rhythms keep different cognitive processes in lock-step, the artificial system could route signals to the right module at the right frequency, avoiding bottlenecks and reducing redundant computation ²⁵ ²⁷. This is a direct parallel to attention routing: the oscillatory scheme acts like an internal traffic signal, regulating when each component should be active and listen to others.

Implications for AI Attention and Decision Routing

Across these domains, a common theme emerges: **phase-based control and routing** can orchestrate complex behavior in a flexible yet robust way. In robotics and RL, phase oscillators coordinate low-level actuators, yielding smooth and adaptive movements; in multi-agent systems, shared oscillation patterns enable agents to **coordinate decisions or actions** without a central clock; in neural architectures, oscillations provide a timing signal that can gate information flow and bind distributed representations (analogous to attention). These examples serve as **templates for general AI design**. For instance, an advanced AI system might assign a unique oscillatory phase or rhythm to different "modules" or agents (such as vision, language, and planning modules) and then regulate their interaction by phase alignment – only modules in sync at a given moment exchange data, preventing interference. Such an **attention routing mechanism** would be dynamic (since phases can continuously shift or re-synchronize) and interpretable (we can literally observe which components are phase-locked and thus communicating). Another implication is improved **robustness and efficiency**: as seen with entrainment in CPGs and synchrony in multi-agent teams, oscillatory coupling lets systems recover from perturbations and coordinate with minimal communication, which is attractive for fault-tolerant AI and distributed learning. Furthermore, leveraging multiple frequencies could allow an AI to handle **multiple scales of decision-making simultaneously** – a slow rhythm could govern strategic planning while faster rhythms handle tactical reactions, all integrated harmoniously. In summary, phase-based techniques inspired by biology are increasingly informing AI **attention and decision-routing strategies**, promising systems that can juggle

many components and tasks by using timing as the organizing principle. The research surveyed here underscores both theoretical foundations and empirical successes that pave the way for future AI that “thinks in rhythms.”

Sources: The information and examples above were drawn from a mix of recent academic literature and implementations. For instance, Sakai *et al.* ³ ²⁸ and Bellegarda & Ijspeert ⁷ ¹ detail how RL can tune CPG oscillators for quadrupedal robots, while Mitra ¹¹ ¹³ uses the Kuramoto model to analyze multi-agent synchronization. Neural and cognitive insights were based on studies like Grothe *et al.* on gamma-band gating in attention ¹⁹ ²⁰, Rajagopal *et al.*’s DONN architecture ²³ ²⁴, and Wang *et al.*’s multi-frequency brain-inspired model ²⁵ ²⁶, among others. These sources provide deeper theoretical background and empirical data on the benefits of oscillator-based control and routing in various AI contexts.

¹ ⁷ ⁸ [2211.00458] CPG-RL: Learning Central Pattern Generators for Quadruped Locomotion

<https://arxiv.labs.arxiv.org/html/2211.00458>

² ³ ⁴ ⁵ ⁶ ²⁸ Hierarchical reinforcement learning with central pattern generator for enabling a quadruped robot simulator to walk on a variety of terrains | Scientific Reports

https://www.nature.com/articles/s41598-025-94163-2?error=cookies_not_supported&code=c36b4911-c31f-4a3fbff-9ecf0886ad85

⁹ ¹⁰ web.mit.edu

<https://web.mit.edu/~jadbabai/www/papers/AntonisUlrichDelay.pdf>

¹¹ ¹² ¹³ ¹⁴ ¹⁵ ¹⁶ [2508.12314] Synchronization Dynamics of Heterogeneous, Collaborative Multi-Agent AI Systems

<https://arxiv.org/abs/2508.12314>

¹⁷ ¹⁸ ²³ ²⁴ Deep oscillatory neural network | Scientific Reports

https://www.nature.com/articles/s41598-025-24837-4?error=cookies_not_supported&code=2e652493-2457-4143-8fd7-cd0bd886e33a

¹⁹ ²⁰ ²¹ Gamma-band synchronization between neurons in the visual cortex is causal for effective information processing and behavior | Nature Communications

https://www.nature.com/articles/s41467-025-62732-8?error=cookies_not_supported&code=19c25382-80b6-488e-8d59-9bbe150b43b1

²² Soft Contrastive Learning for Irregular Multivariate Time Series

<https://openreview.net/pdf?id=9Mx8YfLDbu>

²⁵ ²⁶ ²⁷ Neuromorphic Computing with Multi-Frequency Oscillations: A Bio-Inspired Approach to Artificial Intelligence

<https://arxiv.org/html/2508.02191v1>