



Topology-Driven UX Insights: Persistent Homology in Design

Overview: Uncovering Latent Structure with TDA

Topological Data Analysis (TDA) provides tools like **persistent homology** to reveal hidden geometric structure in complex datasets. Instead of focusing only on pointwise statistics, persistent homology examines how data “points” connect across scales, detecting features such as clusters, loops, and voids that correspond to intrinsic shape. For example, **neural population activity** can lie on low-dimensional manifolds: head-direction cell activity forms a ring (1D loop), and grid-cell ensembles famously form a **toroidal** manifold. In the latter, **persistent homology** confirmed a torus shape by finding Betti numbers $\beta_0=1$, $\beta_1=2$, $\beta_2=1$ ¹ – indicating one connected component, two independent 1D cycles (loops), and one hollow 2D cavity (like the donut hole) ¹. This analysis validated the **latent geometry** of grid cell representations, showing the animal’s positional coding indeed lives on a torus. More broadly, TDA has exposed meaningful shapes in data ranging from neural dynamics to user behavior. For instance, loops in **clickstream data** can signify users repeatedly cycling through pages (a structural pattern invisible to simple funnels). By capturing these shapes, persistent homology uncovers “intrinsic coordinates” of complex systems that can inform design and decision-making.

Topological Metrics and Signatures

Key **metrics from persistent homology** help quantify these latent structures:

- **Betti Numbers ($\beta_{₀$, $\beta_{₁$, $\beta_{₂$, ...}:** Counts of topological features in each dimension. β_0 is the number of connected components, β_1 the number of independent loops or tunnels, β_2 the number of voids or enclosed cavities, etc. For example, a classic torus (donut shape) has $\beta = [1, 2, 1]$ – one component, two distinct loops, one void ¹. A simple loop (circle) has $\beta = [1, 1, 0]$, while a pair of separate loops would have $\beta_0=2$ (two components) unless joined. These **Betti patterns** act as shape signatures. In practice we examine whether the data’s persistent homology yields the expected pattern (e.g. 2 persistent 1D loops and 1 persistent 2D hole for a torus). Consistently observing $\beta = [1,2,1]$ across scales gives confidence of a toroidal structure, whereas deviations (e.g. an extra $\beta_0 > 1$ or missing β_2) suggest fragmentation or collapse of that shape.
- **Persistence Bar Lengths:** Persistent homology doesn’t just count features, it measures their stability across a filtration (e.g. increasing distance threshold in a Vietoris-Rips complex). Each topological feature appears at some scale and disappears at a larger scale, yielding a “barcode” of bars whose length (lifespan) indicates significance. **Long bars** represent robust features (e.g. a long-lived loop likely reflects a true recurring cycle in the data), whereas **short bars** are usually noise. To systematically assess structure, we define **persistence thresholds**: a feature is considered real if its bar length exceeds a chosen cutoff. For example, one might require loops to persist for >10% of the filtration span or beyond a certain distance to be counted. In the grid-cell torus analysis, researchers

applied a common cutoff to classify simulations by Betti numbers; ~82% of trials clearly showed [1,2,1] after thresholding, with the remainder having small deviations consistent with noise ². This approach filters out ephemeral loops, ensuring that only meaningful topological features inform the metrics.

- **“Torus Confidence” (Composite Metric):** We can combine the above indicators into a single **torusConfidence** score that quantifies how strongly the data exhibits a toroidal structure. For example, torusConfidence could be high when $\beta = [1,2,1]$ is observed **and** the two longest β_1 bars plus the β_2 bar far exceed noise thresholds. Concretely, one might define torusConfidence as the minimum persistence (normalized length) of the second-longest 1D hole and the 2D hole – a high minimum implies both loops and the void are well-formed. Additional factors can refine this score: e.g. penalize extra components ($\beta_0 > 1$) or extra loops ($\beta_1 > 2$) as signs of fragmentation. The result is a 0–1 metric (or percentage) indicating certainty that the data’s shape is a single torus rather than disjoint loops or a collapsed structure. A **high torusConfidence** means the data likely lies on a donut-like manifold, whereas lower values may indicate that loops are fleeting or that the topology is simpler (or more complex) than a torus. This metric provides an at-a-glance summary for the UI or system logic to act upon.
- **Persistence Diagram Peaks and Gaps:** Another useful diagnostic is looking at the ranked lifetimes of features. Often, a clear separation (gap) between a few long-lived bars and the rest suggests a well-defined topology. For instance, if exactly two H_1 bars and one H_2 bar tower above all others, that is a strong torus signature. Detecting the **largest gap** in the sorted bar lengths can automatically choose a persistence cutoff (the gap just after the significant features) ². This data-driven thresholding ensures we capture genuine structure even without manually tuning a fixed cutoff. Metrics like the number of bars above the gap or the total “persistent volume” can further quantify topology. For example, one could define a **loop prominence score** as the sum of lengths of the top two H_1 bars (the intuition being that a torus would have two very prominent loops). Such metrics complement Betti counts by incorporating how *pronounced* each feature is.

Using these metrics in tandem provides a robust assessment of topological structure. Betti numbers give the **type** of shape (how many holes of each dimension), while persistence-based measures give the **confidence** in those features. Together they allow a system to distinguish, say, a true torus from a mere pair of unrelated loops or from a degenerate case where loops exist only at a narrow scale. (Notably, advanced invariants like **persistent cup-length** can disambiguate shapes with identical Betti numbers by checking interactions between loops, but in many UX scenarios simpler metrics suffice.) In summary, we track **connectedness (β_0)** to catch fragmentation, **1D loops (β_1)** to see cyclic patterns, **2D voids (β_2)** for toroidal surfaces, and use persistence thresholds to ensure these features are real and significant.

Topology-Driven UI Hooks and Adaptive Logic

By exposing these topological metrics to the interface, we can create **actionable UI hooks** – elements of the user experience that respond to the user’s underlying “manifold state.” Below are concrete ways to tie Betti numbers and related metrics into UI feedback and behavior:

- **Topological Badges (Betti Pattern Indicators):** The UI can display a small badge or icon reflecting the current Betti signature of the user’s data (or interaction pattern). For example, if the system detects $\beta = [1,2,1]$, a **donut-shaped icon** could light up to signify a torus-like pattern (two

independent loops) has formed. If β_1 drops to 0 (no loops), the icon might switch to a simple dot or circle, indicating a loop-free state. Likewise, $\beta_0 > 1$ (multiple disconnected components) could trigger a fractured icon (e.g. split circles or scattered dots) warning that the user's behavior or data is fragmented. These badges act as at-a-glance indicators: e.g. a "donut badge" conveys that the user is in a stable cyclical routine (torus), whereas a "broken-link badge" would signal dispersion. The mapping from Betti numbers to badge visuals should be clear and learnable – perhaps with tooltips (e.g. "Torus mode: two concurrent loops detected") for clarity. This provides continuous ambient feedback about the structure of user activity or state.

- **Live "Manifold Scanner" Feedback:** We can implement a live topology scanner – a UI panel or overlay that updates in near-real-time with the current topological readings. For instance, a small widget could show β_0 , β_1 , β_2 values updating as the user interacts, much like a network signal strength indicator but for shape. The scanner might use simple text or bars (e.g. "Components:1, Loops:2, Voids:1") or even animate an abstract shape morphing according to the detected topology. This gives power-users or researchers immediate insight into how their actions affect the underlying data manifold. In a learning or creative application, the scanner could encourage users to achieve certain structures (e.g. guiding them to form a loop by revisiting earlier steps). The feedback loop can also aid debugging: if a new feature causes user behavior to scatter (β_0 spikes), designers immediately see the topology change and can respond. A **threshold alert** can be built in: for example, if torusConfidence exceeds 0.8, the scanner panel might highlight "Torus detected!" in green; if it falls or if loops disappear ($\beta_1 \rightarrow 0$), it might flash a warning or suggestion.
- **Trigger Logic for Adaptive UI States:** The system can use topology changes as triggers to adapt the UI or workflow in real-time. Because topological features often correspond to user states (e.g. being "stuck in a loop" or exploring new territory), reacting to them can improve UX. For example, if **β_1 drops** from 1 to 0 (meaning a loop has just broken), it could indicate the user resolved a previously repetitive cycle – the UI might then advance to a new tutorial step or increase difficulty, capitalizing on the breakthrough. Conversely, if **β_1 rises** or stays high (the user keeps looping through certain actions), the system might detect a potential **looping pattern**. This could trigger an intervention: a contextual hint, an offer of help ("It looks like you're revisiting the same steps – need assistance?"), or a gentle nudge to break the cycle. Research suggests a persistent loop in user navigation can indicate circling without progress – a chance to optimize UX or prompt the user differently. Similarly, a **rise in β_0** (fragmentation) means the user's session has split into disjoint parts (maybe multi-tasking or confusion causing divergence). The application could respond by introducing a unifying guide (e.g. "Let's refocus on one task") or merging interface elements to encourage convergence. These triggers should have some hysteresis (to avoid flapping): for instance, require $\beta_0 > 1$ for a sustained period (several seconds or interactions) before declaring "fragmentation" and altering the UI. By tying UI logic to the **manifold state**, we create an adaptive system that responds to deeper patterns in user behavior rather than just surface events.
- **Adaptive Content and Difficulty:** Topological feedback can also inform content generation or difficulty tuning on the fly. Imagine a learning app that tracks a student's problem-solving path as a point cloud; a loop might mean the student is cycling between a few concepts repeatedly. If torusConfidence gets high (indicating two independent loops, perhaps oscillating between two problem types), the app might introduce an integrative task to connect those loops (an attempt to "untangle the donut"). A collapse of loops ($\beta_1 \rightarrow 0$ with $\beta_0=1$) could indicate the student has converged on a solution path, so the app can safely increase difficulty or move to the next module.

These state-based adjustments, triggered by topological metrics, ensure the interface remains in sync with the user's cognitive state.

To implement these hooks robustly, it's crucial to **smooth and debounce** the triggers (as discussed in the next section). The UI should not jitter or flip states on every minor metric change – instead, we watch for significant, sustained topological shifts. Overall, by linking Betti numbers and persistence-based insights to UI elements, we enrich the interface with an awareness of the user's journey "shape." This makes the UX more responsive: the system can detect when users are in a dead-end loop, wandering in multiple threads, or cruising on a stable torus, and adjust the experience proactively.

Implementing a Real-Time Persistent Homology Pipeline

Building a lightweight persistent homology pipeline for real-time UX contexts requires balancing **speed, stability, and accuracy**. Below we outline practical guidance for implementation:

- **Data Windowing & Downsampling:** Rather than using an ever-growing dataset, operate on a sliding window of recent user data or states. For example, consider only the last N user interactions or last T minutes of activity to construct the point cloud or time series for analysis. This bounds the problem size and focuses on the current state. Within that window, downsample aggressively: you might randomly sample or use **landmark points** to reduce complexity while preserving shape. Landmark-based Vietoris-Rips filtrations can approximate the full topology with far fewer points. For instance, selecting ~100–200 landmark points (out of potentially thousands of recent data points) can capture the main loops/holes with small error. This is justified by stability theorems – a coarse point cloud can still reflect the correct homology of the underlying manifold if sampling is sufficient. Utilizing landmarks or clustering (e.g. cluster points and use cluster centroids) will dramatically cut computation costs.
- **Efficient PH Computation:** Use optimized libraries and limit dimensions. Fast implementations like **Ripser** (in Python via Ripser.py) or libraries like **Giotto-TDA** and **Scikit-TDA** are well-suited for computing persistent homology on the fly. These libraries can compute H0, H1 (and H2 if needed) efficiently on moderately sized point clouds (hundreds of points). Limit the homology dimensions to what you need – for detecting loops and a torus, you may only need up to H2. Skipping higher dimensions saves time. Also consider simpler filtrations: if the data is inherently low-dimensional (say you always embed user state in 3D via PCA), using an **alpha complex** (Delaunay-based) or even just computing a *distance graph* for H1 might suffice. The goal is to minimize overhead so that updates can be frequent (e.g. every second or in response to each batch of new interactions).
- **Persistence Thresholds & Feature Extraction:** Decide on thresholds or criteria to extract features from the persistence output automatically. One approach is to pick a fixed **ϵ -threshold** that defines the scale at which you read off Betti numbers (essentially slicing the barcode at a particular filtration value). However, a more robust approach is to use the **lifetime gap heuristic** discussed earlier: identify a large gap in the sorted bar lengths for each homology dimension and cut there ². For instance, if two H1 bars are much longer than the rest, we classify $\beta_1 = 2$. This way the system adapts to the data's scale and noise level. Implement this by sorting persistence intervals after each PH computation and checking length differences. The output of this step should be the **current Betti numbers** (for significant features) and perhaps the normalized lengths of those features. These can feed directly into UI logic (e.g. set a boolean `hasLoop = True` if a loop bar > 0.1 persists).

- **Smoothing and Debouncing:** Because each recomputation of homology might fluctuate (especially if the input window is shifting), it's important to smooth the metrics to avoid a flickery UI. Maintain a short moving average or exponential smoothing on key metrics like torusConfidence or the counts of loops. For example, you could average the last 5 values of β_1 or use a time-decay smoothing so that transient blips are damped. Additionally, implement **hysteresis** for state changes: require that a condition is met continuously for a certain duration before acting. If β_1 has dropped to 0 but only for one frame, don't immediately declare the loop gone; wait until it remains 0 for say 3 consecutive scans (or 2 seconds of time) before triggering the UI change. Similarly, when a new loop appears, perhaps wait for it to persist beyond a minimal lifetime threshold. This prevents the UI from reacting to noise or momentary partial structures. In essence, treat the topology stream as you would a noisy sensor signal – apply low-pass filtering and change detection with confirmation.
- **Performance and Frequency:** Decide the update frequency based on user context and performance budgets. If analyzing a rapid stream (e.g. user EEG or fine-grained cursor trajectories), you might recompute PH every second or continuous background thread. For slower interactions (page views, clicks), on-demand computation when a session event occurs may suffice. Profile the pipeline: computing persistence on 100 points up to H1 is extremely fast (tens of milliseconds) with modern algorithms; including H2 or more points might take longer but can often be kept under a second with optimizations. If needed, use asynchronous processing – compute topology on a separate thread or worker so as not to freeze the UI. Should performance still lag, consider more drastic simplifications: e.g. **pre-compute** a lookup table of typical patterns, or use an online learning model trained to predict Betti numbers from simpler streaming features (though that somewhat defeats the purpose of rigorous TDA). Usually, a well-chosen combination of downsampling and efficient libraries will keep the pipeline lightweight. For example, researchers analyzing neural manifolds have successfully used ~150 landmarks and integration with Ripser to handle large datasets in reasonable time. In a UX setting, the data is often smaller scale (hundreds or thousands of user states), so real-time is achievable.
- **Robustness Considerations:** Ensure the pipeline handles edge cases. If the window has too few points (e.g. user just started, or data is sparse), the homology might be trivial or unstable. The code should handle these gracefully (perhaps by defaulting to $\beta_0=1$, $\beta_1=0$ in very low data scenarios, and treating torusConfidence as 0). Monitor memory usage if the window slides – remove old points as new ones come (a fixed-size buffer of points can help). Use numeric stability tricks from TDA libraries, such as working with approximate distances or adding a tiny noise to avoid degeneracies (no two distances exactly equal, etc.). Finally, test the pipeline with known synthetic shapes (generate a torus point cloud, a random loop, noise, etc.) to verify that your thresholds and confidence metrics correctly identify shapes when they should, and yield low scores when no shape is present. Tuning on synthetic "ground truth" shapes can calibrate the thresholds for the real usage context.

By following these implementation practices, you can embed a **persistent homology engine** into an interactive system that streams topology insights without a heavy performance penalty. The combination of windowed data, smart downsampling, optimized PH computation, and smoothed output will produce stable topological metrics ready for driving UI changes.

Visualizing Topology in the Interface

Finally, we consider **visual design concepts** to represent topological insights in an intuitive and engaging way. The goal is to surface metrics like Betti numbers, loops, and holes through visuals that make sense to end-users (or analysts) at a glance. Below are proposed UI elements and plots to communicate topology:

- **Betti Badges and Icons:** As mentioned, using icons can compress complex topological states into a simple visual. A **donut icon** is a natural choice for indicating a torus (since a torus is literally a donut shape in topology). This could be stylized as a torus badge that appears in the UI when a toroidal structure is detected with high confidence. The icon might even be annotated with small numbers or rings to indicate the loops: e.g. a donut with two distinct colored rings around it to signify $\beta_1=2$. For a single-loop state, a simple ring icon (like a circle or loop arrow) can be used. For no loops ($\beta_1=0$), the icon might be filled in (solid disk) showing nothing “hollow” inside. And if multiple components appear ($\beta_0>1$), the icon could split into multiple pieces or an emoji-like symbol of scattered nodes. The use of familiar shapes (donut = torus, circle = loop, cluster of dots = components) leverages the intuitive side of topology. By clicking or hovering on these badges, users could get a tooltip like “Current state: 2 loops detected (toroidal)” for further explanation. These badges make topology **visible** and friendly, rather than an abstract math concept.
- **Mini Persistence Bar Charts:** To give a more quantitative feel, the UI can include a tiny **barcode chart** or histogram indicating the persistence of topological features. For example, a small widget could display bars labeled “Loop1”, “Loop2”, “Void” corresponding to the lengths of the longest one or two H1 bars and the H2 bar. The bars could fill up relative to a max scale (perhaps the current threshold or the maximum possible persistence in the data). If a feature is absent or below threshold, its bar might be empty or a different color. This mini chart allows users to see not just that loops exist, but how strong they are. It could be represented as a stacked icon: imagine a donut icon with two radial progress rings around it – each ring’s completion represents the persistence of one loop relative to full confidence. Another design: a tiny barcode image where long lines indicate strong features and short lines indicate weak ones (with maybe noise lines faded out). By presenting this, the interface gives a peek into the *barcodes* that underlie the Betti counts, which can build user trust (they can see a second loop is barely there versus clearly persistent). It’s like a “signal strength” indicator but for each topological feature.
- **Torus Wireframe Overlays:** In scenarios where you visualize user data or state in a spatial layout (e.g. embedding user behavior in 2D/3D space via PCA or t-SNE), overlaying a reference shape can be powerful. If the system detects a torus, the UI can draw a faint **wireframe torus** in the background of the scatter plot, orienting it to align with the detected loops. This helps users literally see the donut that their data points lie on. For instance, a 3D scatterplot of the user’s state trajectory could suddenly reveal a donut structure when the wireframe is superimposed, making the abstract Betti numbers concrete. Similarly, if a single loop is detected ($\beta_1=1$), a circle or ring curve could be drawn through the data points, suggesting the cyclical path the user is taking. These overlays should be subtle (perhaps a dashed grid) and possibly interactive (the user could toggle them or rotate the view). They act as visual **affordances** for complex geometry – turning invisible manifolds into ghost-like guides in the interface. In a more stylized UI (say a dashboard for monitoring user behavior), the torus wireframe could even be an iconographic element that appears next to user segments or profiles that exhibit toroidal patterns.

- **Phase-Space Trajectory Plots:** To reflect **user state cycles** or interaction loops, a dedicated phase plot can be used. This could be a small chart (like a sparkline) that plots the user's path through a reduced state space over time. For example, using two principal components that capture the dominant periodic behaviors, we could plot the user's trajectory as a line that often loops around. If the user has two independent cyclical behaviors (forming a torus in higher dimensions), the phase plot might show a Lissajous-like figure or a dense loop that shifts over time. The UI can highlight loops by detecting when the trajectory line closes on itself (approximately) and perhaps coloring that segment. If the user's attention or activity goes through cycles (e.g. alternating between two modes), the phase plot will make this visible as a repeating orbit. One design idea is an "**attention orbit**" display: a circle graph where the angle represents one aspect of user state and radius another, so that consistent looping behavior draws a stable shape (circle or figure-eight) on it. When a toroidal pattern is present, the combination of two frequencies might produce a donut-like shape when visualized appropriately (e.g. a rotating 3D phase plot). While these plots are more advanced, they provide a **rich visualization** for power users or researchers to diagnose *how* the user is looping. Importantly, they tie back to the real meaning: a loop in the plot corresponds to the user revisiting states in an **almost periodic** way (perhaps switching back-and-forth between two pages regularly, etc.), which is exactly what $\beta_1 > 0$ indicated. Thus, the visualization closes the loop (pun intended) between the quantitative metric and the user's lived experience in the interface.
- **Color and Animation to Emphasize Changes:** In all the above visual elements, using color-coding or animation can help draw attention to changes in topology. For instance, the Betti badge icon might be gray when no significant topology is present, but turn vibrant (or start spinning slowly like a wheel) when a loop emerges. The persistence bar chart could gently pulse when a bar grows longer (indicating a feature solidifying). If a torus wireframe overlay is active, it might fade in or glow softly when torusConfidence passes a high threshold, reinforcing that a strong structure has been found. Conversely, if the structure breaks (say one of the two loops disappears), the wireframe could crack or the donut icon could show a break, providing immediate visual indication that something changed. These design touches ensure that topological feedback isn't overlooked – the UI should treat a major topology change as a notable event (just as it would a critical alert). By making topology **perceptible and dynamic**, users and designers can incorporate these insights naturally into their workflow.

In summary, the visual design should translate algebraic topology into intuitive metaphors: donuts, loops, holes, and surfaces that users can see and relate to their context. Prior art in topology often uses playful analogies (coffee cups and donuts); we extend that to an interactive setting with icons and overlays that truly bring the "shape of data" to life. Together with the earlier metrics and pipeline, these visualization concepts form a complete loop: data → homology → metrics → UI representation. This empowers interface designs that are **shape-aware**, reacting not just to what users do, but to the underlying patterns in *how* they do it – closing the gap between data science insights and user experience.

References and Applications

By grounding these concepts in current research, we ensure our approach is credible and cutting-edge. The use of persistent homology to find a torus in grid-cell activity is a **validated neuroscience result**¹, showcasing TDA's power in revealing continuous attractor manifolds. Similarly, TDA has been applied to behavior and cognitive data: for example, analysis of web navigation paths has used persistent homology to uncover loops and voids in user journeys. This multiscale view revealed "intricate patterns and structures

that other methods might overlook,” highlighting TDA’s value in UX research. In practical analytics, recent work notes that **long persistent loops often correspond to meaningful user behavior cycles**, such as customers “looping” between product pages and carts, which signal UX pain points. By leveraging these insights, our proposed design marries proven TDA techniques with UI innovation. The result is an interface that is aware of the *topology* of user interaction, potentially leading to new forms of adaptive, intelligent UX. This approach is at the frontier of data-driven design, demonstrating how concepts like Betti numbers and persistent homology can move from pure analysis into real-time user-facing features – effectively turning the “shape of data” into a tangible part of the user experience.

<Rigidbody style="display:none;"> 1 2 </Rigidbody>

1 2 (PDF) Modeled grid cells aligned by a flexible attractor

https://www.researchgate.net/publication/386487524_Modeled_grid_cells_aligned_by_a_flexible_attractor